

```
In [1]: import pandas as pd
import numpy as np
```

```
In [2]: df = pd.read_csv('height-weight.csv')
```

```
In [3]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 23 entries, 0 to 22
Data columns (total 2 columns):
 #   Column  Non-Null Count  Dtype  
---  -
 0   Weight  23 non-null      int64  
 1   Height  23 non-null      int64  
dtypes: int64(2)
memory usage: 496.0 bytes
```

```
In [4]: df.head()
```


```
Out[4]:
```

	Weight	Height
0	45	120
1	58	135
2	48	123
3	60	145
4	70	160

```
In [7]: import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
In [6]: plt.scatter(df.Weight,df.Height)#the data is following a linear regression
plt.xlabel('Weight')
plt.ylabel('Height')
```

```
Out[6]: Text(0, 0.5, 'Height')
```




```
In [8]: df.corr()#weight and height are positively and strongly correlated
```

```
Out[8]:
```

	Weight	Height
Weight	1.000000	0.931142
Height	0.931142	1.000000

```
In [12]: sns.pairplot(df)
```

```
Out[12]: <seaborn.axisgrid.PairGrid at 0x7fa32ac54fa0>
```



```
In [13]: X = df[['Weight']]#assigning independant variable to X
#since X should be a dataframe we have used [][] two brackets
```

```
In [18]: X.shape
```

```
Out[18]: (23, 1)
```

```
In [19]: Y = df['Height']#assigning dependat variable to Y. Here we have used [] as Y can be a se
```

```
In [20]: Y.shape
```

```
Out[20]: (23,)
```

```
In [22]: from sklearn.model_selection import train_test_split
```

```
In [23]: X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size= 0.3, random_state=
```

```
In [24]: from sklearn.preprocessing import StandardScaler
```

```
In [26]: scaler = StandardScaler()
```

```
In [28]: X_train = scaler.fit_transform(X_train)
```

```
In [29]: X_test = scaler.transform(X_test)#using tranform only instead of fit_transform to avoid
```

```
In [30]: from sklearn.linear_model import LinearRegression
```

```
In [31]: regression = LinearRegression()
```

```
In [32]: model = regression.fit(X_train, y_train)#building regression model
```

```
In [33]: model.coef_
```

```
Out[33]: array([17.03207732])
```

```
In [34]: print('Coefficient or Slope:', model.coef_)
```

```
Coefficient or Slope: [17.03207732]
```

```
In [35]: model.intercept_
```


```
Out[35]: 157.8125
```

```
In [36]: print('Intercept:', model.intercept_)
```

```
Intercept: 157.8125
```

```
In [37]: plt.scatter(X_train, y_train)  
plt.plot(X_train, model.predict(X_train))#plotting best fit lineb
```

```
Out[37]: [<matplotlib.lines.Line2D at 0x7fa30ef2f040>]
```



```
In [38]: #performing prediction  
y_pred = model.predict(X_test)
```

```
In [39]: y_pred
```

```
Out[39]: array([162.55745791, 162.55745791, 128.70154204, 179.99838426,  
149.22027893, 190.25775271, 142.03872102])
```

```
In [40]: #finding performance  
from sklearn.metrics import mean_squared_error, mean_absolute_error
```

```
In [41]: mse = mean_squared_error(y_test, y_pred)
```

```
In [42]: mse
Out[42]: 103.09818711844574

In [43]: mae = mean_absolute_error(y_test, y_pred)
mae
Out[43]: 9.237776679921925

In [44]: rmse = np.sqrt(mse)
rmse
Out[44]: 10.153727744944009

In [45]: from sklearn.metrics import r2_score

In [47]: r2score = r2_score(y_test, y_pred)
r2score #shows the model performs at 78%
Out[47]: 0.7828485570493535

In [50]: #prediction for new data that is predicting height by entering weight
model.predict(scaler.transform([[72]]))#here the height is 156cm if the weight is 72kg

/Users/ishutejwani/opt/anaconda3/lib/python3.9/site-packages/sklearn/base.py:465: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names
  warnings.warn(
Out[50]: array([156.40183684])

In [51]: #performing regression via OLS
import statsmodels.api as sm

In [52]: olsmodel = sm.OLS(y_train, X_train).fit()

In [53]: olsmodel.summary()# the coefficient here is almost same as the coefficient arrived earlier

/Users/ishutejwani/opt/anaconda3/lib/python3.9/site-packages/scipy/stats/_stats_py.py:1769: UserWarning: kurtosistest only valid for n>=20 ... continuing anyway, n=16
  warnings.warn("kurtosistest only valid for n>=20 ... continuing ")
Out[53]:
```

OLS Regression Results

Dep. Variable:	Height	R-squared (uncentered):	0.012
Model:	OLS	Adj. R-squared (uncentered):	-0.054
Method:	Least Squares	F-statistic:	0.1745
Date:	Sat, 18 Nov 2023	Prob (F-statistic):	0.682
Time:	14:06:19	Log-Likelihood:	-103.69
No. Observations:	16	AIC:	209.4
Df Residuals:	15	BIC:	210.2
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
x1	17.0321	40.767	0.418	0.682	-69.861	103.925

Omnibus:	0.957	Durbin-Watson:	0.003
Prob(Omnibus):	0.620	Jarque-Bera (JB):	0.320
Skew:	-0.347	Prob(JB):	0.852
Kurtosis:	2.998	Cond. No.	1.00

Notes:

[1] R^2 is computed without centering (uncentered) since the model does not contain a constant.

[2] Standard Errors assume that the covariance matrix of the errors is correctly specified.