

```
In [1]: import pandas as pd
import numpy as np
```

```
In [2]: indexprice_df = pd.read_csv('economic_index.csv')
```

```
In [3]: indexprice_df.info() #there are no null values as all columns have 24 records
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 24 entries, 0 to 23
Data columns (total 6 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   Unnamed: 0             24 non-null    int64
1   year                  24 non-null    int64
2   month                 24 non-null    int64
3   interest_rate          24 non-null    float64
4   unemployment_rate      24 non-null    float64
5   index_price            24 non-null    int64
dtypes: float64(2), int64(4)
memory usage: 1.2 KB
```

```
In [4]: indexprice_df.head() #we require interest_rate and unemployment_rate to predict the index price
#We can remove unwanted columns
```

```
Out[4]:
```

	Unnamed: 0	year	month	interest_rate	unemployment_rate	index_price
0	0	2017	12	2.75	5.3	1464
1	1	2017	11	2.50	5.3	1394
2	2	2017	10	2.50	5.3	1357
3	3	2017	9	2.50	5.3	1293
4	4	2017	8	2.50	5.4	1256

```
In [6]: indexprice_df.drop(columns= ['Unnamed: 0', 'year', 'month'], axis= 1, inplace= True)
```

```
In [7]: indexprice_df.head()
```


```
Out[7]:
```

	interest_rate	unemployment_rate	index_price
0	2.75	5.3	1464
1	2.50	5.3	1394
2	2.50	5.3	1357
3	2.50	5.3	1293
4	2.50	5.4	1256

```
In [8]: import matplotlib.pyplot as plt
import seaborn as sn
%matplotlib inline
```


```
In [9]: plt.scatter(indexprice_df.interest_rate, indexprice_df.index_price)
plt.xlabel('Interest Rate')
plt.ylabel('Index Price')
#looks like Interest Rate and Index Price are linearly related
```

```
Out[9]: Text(0, 0.5, 'Index Price')
```



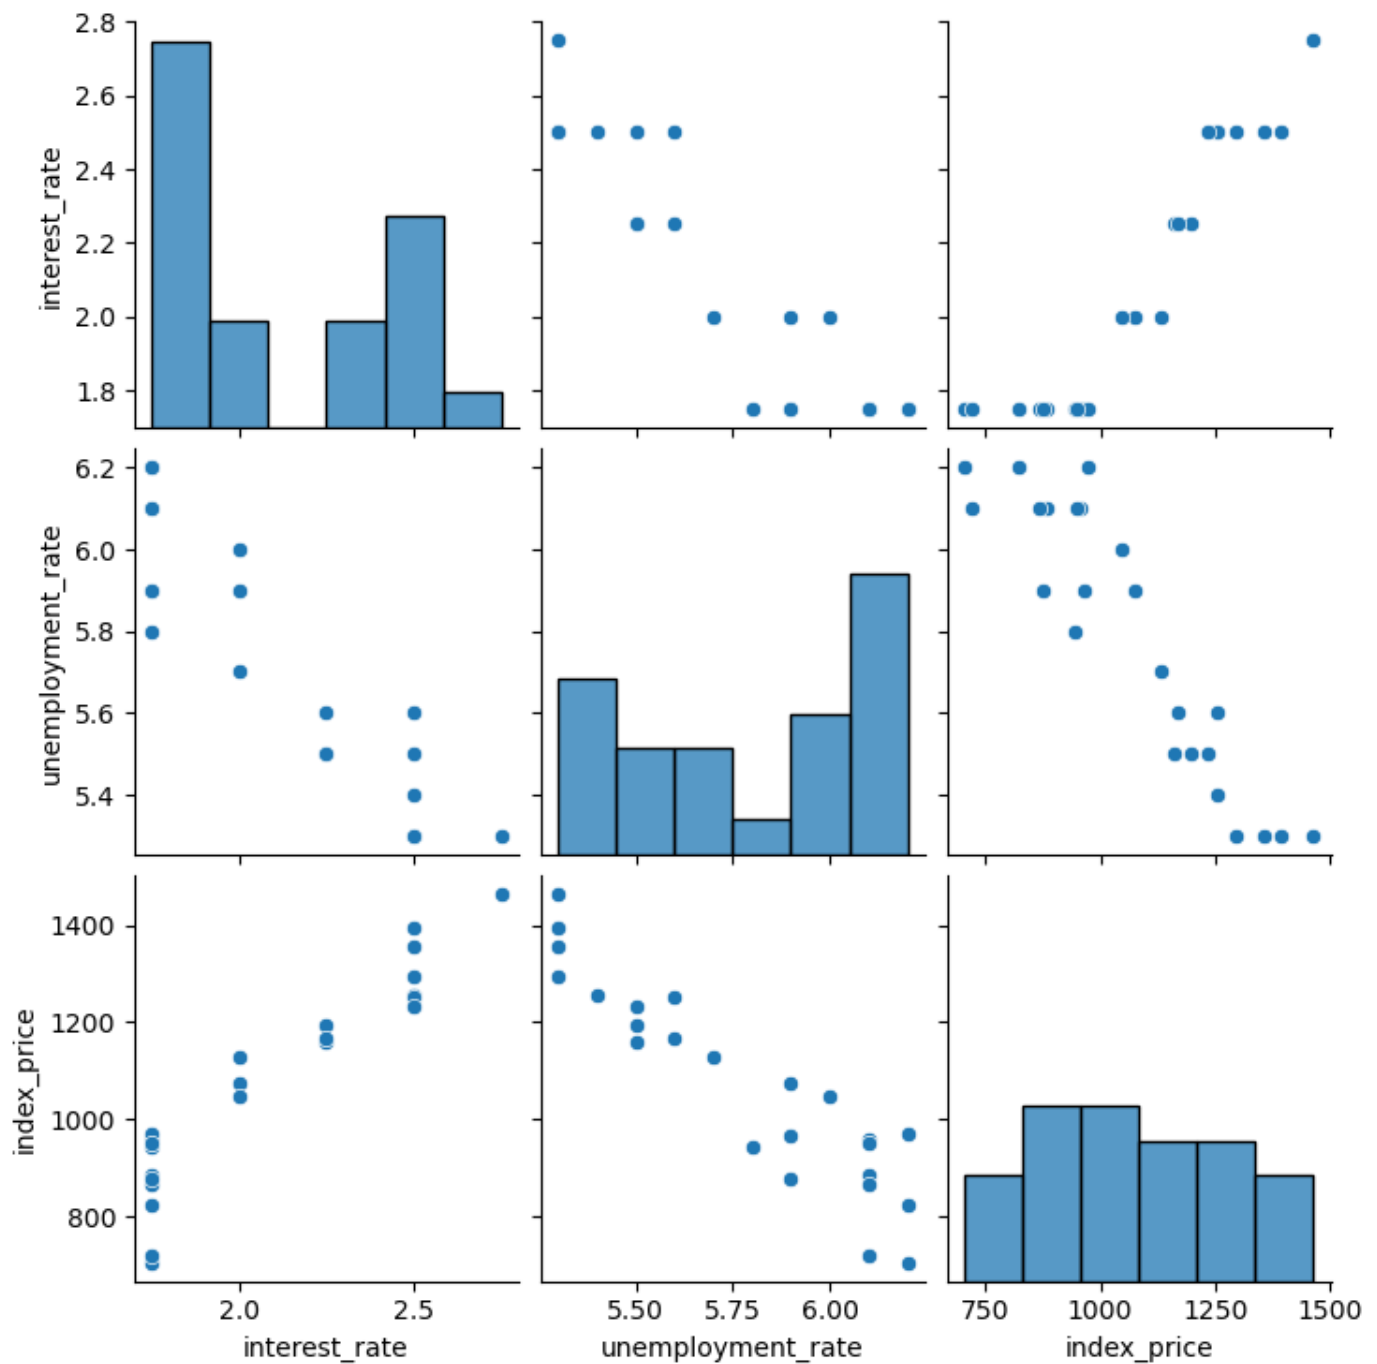
```
In [10]: plt.scatter(indexprice_df.unemployment_rate, indexprice_df.index_price)
plt.xlabel('Unemployment Rate')
plt.ylabel('Index Price')
#looks like unemployment and index price are inversely related
```

```
Out[10]: Text(0, 0.5, 'Index Price')
```



```
In [11]: sn.pairplot(indexprice_df)
#looks like interest rate and unemployment rate are also inversely related
```


Out[11]: <seaborn.axisgrid.PairGrid at 0x7fb20340ed60>



```
In [27]: sn.regplot(indexprice_df.interest_rate, indexprice_df.index_price)
```

```
/Users/ishutejwani/opt/anaconda3/lib/python3.9/site-packages/seaborn/_decorators.py:36:
FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, th
e only valid positional argument will be `data`, and passing other arguments without an
explicit keyword will result in an error or misinterpretation.
warnings.warn(
```

```
Out[27]: <AxesSubplot:xlabel='interest_rate', ylabel='index_price'>
```



```
In [12]: x = indexprice_df.iloc[:, :-1]
x
```

```
Out[12]:
```

	interest_rate	unemployment_rate
0	2.75	5.3
1	2.50	5.3
2	2.50	5.3
3	2.50	5.3
4	2.50	5.4
5	2.50	5.6
6	2.50	5.5
7	2.25	5.5
8	2.25	5.5
9	2.25	5.6
10	2.00	5.7
11	2.00	5.9
12	2.00	6.0
13	1.75	5.9
14	1.75	5.8
15	1.75	6.1
16	1.75	6.2
17	1.75	6.1
18	1.75	6.1

19	1.75	6.1
20	1.75	5.9
21	1.75	6.2
22	1.75	6.2
23	1.75	6.1

```
In [13]: type(X) #independant variable is in dataframe format
```

```
Out[13]: pandas.core.frame.DataFrame
```

```
In [14]: Y = indexprice_df.iloc[:, -1]
Y
```

```
Out[14]: 0      1464
1      1394
2      1357
3      1293
4      1256
5      1254
6      1234
7      1195
8      1159
9      1167
10     1130
11     1075
12     1047
13      965
14      943
15      958
16      971
17      949
18      884
19      866
20      876
21      822
22      704
23      719
Name: index_price, dtype: int64
```

```
In [15]: type(Y) #dependatn variable is in Series format
```

```
Out[15]: pandas.core.series.Series
```

```
In [16]: from sklearn.model_selection import train_test_split
```

```
In [17]: #splitting the data into train and test
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size= 0.3, random_state=
```

```
In [18]: from sklearn.preprocessing import StandardScaler
```

```
In [19]: scaler = StandardScaler()
```

```
In [20]: #standardizing the data
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

```
In [21]: from sklearn.linear_model import LinearRegression
```

```

In [22]: regression = LinearRegression()

In [23]: model = regression.fit(X_train, y_train)

In [24]: model.coef_
Out[24]: array([ 96.28689501, -101.57024663])

In [25]: print('Coefficients or Slope', model.coef_)
Coefficients or Slope [ 96.28689501 -101.57024663]

In [26]: print('Intercept', model.intercept_)
Intercept 1037.6875

In [30]: #performing cross validation
from sklearn.model_selection import cross_val_score
validation_score = cross_val_score(regression,X_train, y_train, scoring='neg_mean_square

In [31]: validation_score
Out[31]: array([-5495.88007512, -7496.90768619, -9203.37178948, -9743.00656351,
               -523.25201601])

In [32]: np.mean(validation_score)
Out[32]: -6492.48362606083

In [34]: y_pred = model.predict(X_test)

In [35]: y_pred
Out[35]: array([1192.13083729,  824.23971817, 1400.41971162,  856.16016713,
               992.22505325, 1160.21038833,  920.00106505, 1328.19572341])


In [37]: from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score

In [38]: mse = mean_squared_error(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)
rmse = np.sqrt(mse)
score = r2_score(y_test, y_pred)
print('Mean Squared Error:', mse)
print('Mean Absolute Error:', mae)
print('Root Mean Squared Error:', rmse)
print('R squared:', score)

Mean Squared Error: 5088.329958294
Mean Absolute Error: 58.95987629034906
Root Mean Squared Error: 71.33253085580239
R squared: 0.8640024299625207

In [39]: plt.scatter(y_test,y_pred)#the data shows that the prediction is linear with absolute d
#the model has performed fairly well
Out[39]: <matplotlib.collections.PathCollection at 0x7fb207a56a00>

```




```
In [40]: residuals = y_test - y_pred
residuals
```

```
Out[40]: 8      -33.130837
16      146.760282
0        63.580288
18       27.839833
11       82.774947
9         6.789612
13       44.998935
1        65.804277
Name: index_price, dtype: float64
```


```
In [41]: sn.displot(residuals, kind='kde')#residuals are normally distributed. this also means t
#performed well
```

```
Out[41]: <seaborn.axisgrid.FacetGrid at 0x7fb2046e7850>
```



In [42]: `plt.scatter(y_pred, residuals)` *#the data is uniformly distributed which means there is no #is performing well*

Out[42]: `<matplotlib.collections.PathCollection at 0x7fb20661fee0>`



In [44]: `#performing prediction by entering interest rate as 3.5 and unemployment rate as 7.5. model.predict(scaler.transform([[3.5, 7.5]]))` *#the predicted index price is 914.84*



```
/Users/ishutejwani/opt/anaconda3/lib/python3.9/site-packages/sklearn/base.py:465: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names
```

```
warnings.warn(
```

```
Out[44]: array([914.84179912])
```